

Dragonfruit AI Challenge

1. Efficient Data Structures for Image Representation

For storing and processing the images, we need efficient data structures that can handle the high resolution (100,000 x 100,000 pixels) while minimizing storage space.

Microscope Images

We are given that microscope images typically have large contiguous regions of black pixels (background) and a single contiguous blob of white pixels (parasite).

Using a **Run-Length Encoding** (RLE) is an efficient way to store images like these.

RLE then compresses the image by storing the lengths of runs of consecutive pixels with the same value.

RLE Representation for Microscope Image:

- Store the pixel value (black (0) or white (1)) and the length of the run.
- Example: [(0, 500), (1, 2500), (0, 97500)] for a row with 500 black pixels, 2500 white pixels, and 97500 black pixels.

Estimated Storage Size:

- In the worst case (an image with alternating black and white pixels), the storage size would be around **100,000 * 2** entries (each run represented by a pair of values), resulting in 200,000 values.
- If each value takes 4 bytes (assuming integer representation), the total size would be 800,000 bytes per row, leading to approximately **80 GB** for the entire image.

We also know that Dye Sensor Images can have more scattered lit areas. Therefore, using a **Sparse Matrix Representation** can be more efficient. We will store only the coordinates of the lit pixels (where dye is present).

Sparse Matrix Representation for Dye Sensor Image:

- Store the coordinates of lit pixels.
- Example: [(x1, y1), (x2, y2), ...] where each tuple (x,y) represents the position of a lit pixel.

Estimated Storage Size:

- Assuming fewer than 10% of pixels are lit, this results in at the very most **10,000,000** coordinates.
- If each coordinate pair takes 8 bytes (4 bytes each for x and y), the total size would be 80,000,000 bytes or **80 MB** for the entire image.

2. Generating Simulated Images

We can simulate images using the above data structures in python:

generate.py:

import random

```
def generate_microscope_image(width, height):
    image = []
    for _ in range(height):
        row = []
        for _ in range(width):
            if random.random() < 0.75: # 75% chance for background (black pixel)
                row.append(0)
            else: # 25% chance for blob (white pixel)
                row.append(1)
        image.append(row)
    return image
```

```
def generate_dye_image(width, height):
    image = []
    for _ in range(height):
        row = []
        for _ in range(width):
            if random.random() < 0.05: # 5% chance for lit pixel (dye present)
                row.append(1)
            else:
                row.append(0)
        image.append(row)
    return image
```

```
if __name__ == "__main__":
    width, height = 100, 100 # Using smaller dimensions for testing
    microscope_image = generate_microscope_image(width, height)
    dye_image = generate_dye_image(width, height)

    # Save generated images to files
    with open("microscope_image.txt", "w") as f:
```

```
for row in microscope_image:
    f.write(','.join(map(str, row)) + "\n")

with open("dye_image.txt", "w") as f:
    for row in dye_image:
        f.write(','.join(map(str, row)) + "\n")

print("Images generated and saved.")
```

Explanation:

generate_microscope_image:

- This creates an image with a black background and a white blob, represented as a 2D list. Each pixel in the list is either 0 (black, representing the background) or 1 (white, representing the blob).

generate_dye_image:

- This creates an image with randomly scattered lit pixels, represented as a 2D list. Each pixel is either 1 (dye present) or 0 (no dye).

Main:

- Uses the generate_microscope_image and generate_dye_image functions to generate two images.
- Saves generated images to two text files (microscope_image.txt and dye_image.txt). Each row of the image is now saved as a comma-separated string of pixel values.

3. Detecting Cancer in Parasites

This function calculates whether the amount of dye within the parasite exceeds 10% of its total area.

detect_cancer.py:

```
def has_cancer(microscope_image, dye_image, width, height):
    total_blob_pixels = 0
    total_dye_pixels = 0

    for y in range(height):
        for x in range(width):
            if microscope_image[y][x] == 1: # Blob pixel
                total_blob_pixels += 1
            if dye_image[y][x] == 1: # Dye within blob
                total_dye_pixels += 1

    dye_percentage = total_dye_pixels / total_blob_pixels if total_blob_pixels > 0 else 0
    return dye_percentage > 0.10

if __name__ == "__main__":
    # Load generated images from files
    with open("microscope_image.txt", "r") as f:
        microscope_image = [list(map(int, line.strip().split(','))) for line in f]

    with open("dye_image.txt", "r") as f:
        dye_image = [list(map(int, line.strip().split(','))) for line in f]

    width, height = len(microscope_image[0]), len(microscope_image) # Get dimensions from image
    result = has_cancer(microscope_image, dye_image, width, height)
    print(f"Parasite has cancer: {result}")
```

Explanation:

- The has_cancer function reads the images from the files, splitting each line by commas to reconstruct the 2D lists.
- The dimensions of the images are determined by the length of the lists.
- The function checks for total_blob_pixels > 0 to avoid division by zero.
- Main loads the generated picture from the file

4. Optimizing Execution Speed

We optimize the execution speed using NumPy

optimize.py:

```
import numpy as np

def has_cancer_optimized(microscope_image, dye_image, width, height):
    # Convert lists to NumPy arrays for efficient processing
    microscope_array = np.array(microscope_image)
    dye_array = np.array(dye_image)

    # Count the number of blob pixels and dye pixels within the blob
    blob_pixels = np.sum(microscope_array == 1)
    dye_pixels_within_blob = np.sum((microscope_array == 1) & (dye_array == 1))

    # Calculate the percentage of dye pixels within the blob
    dye_percentage = dye_pixels_within_blob / blob_pixels if blob_pixels > 0 else 0
    return dye_percentage > 0.10

if __name__ == "__main__":
    # Load generated images from files
    with open("microscope_image.txt", "r") as f:
        microscope_image = [list(map(int, line.strip().split(','))) for line in f]

    with open("dye_image.txt", "r") as f:
        dye_image = [list(map(int, line.strip().split(','))) for line in f]

    width, height = len(microscope_image[0]), len(microscope_image) # Get dimensions from image
    result = has_cancer_optimized(microscope_image, dye_image, width, height)
    print(f"Parasite has cancer: {result}")
```

Explanation: We pretty much do the same thing as detect_cancer but with NumPy arrays

- Converts the 2D lists of the microscope and dye images to NumPy arrays.
- Counts the number of blob pixels (parasite body) in the microscope image.
- Counts the number of dye pixels within the blob area using a logical AND operation between the two arrays.
- Calculates the percentage of dye pixels within the blob.
- Returns True if the dye percentage exceeds 10%, indicating cancer; otherwise, returns False.

5. Other Compression Techniques and Their Impact

Compression Techniques:

1. Block-Based Compression:

- **Method:** Divide the image into smaller blocks and compress each block individually.
- **Impact:** Reduces storage space by exploiting redundancy within each block but may introduce slight decompression overhead.

2. Wavelet Transform:

- **Method:** Apply wavelet transformation to compress the image by removing less significant details.
- **Impact:** Significantly reduces storage space while maintaining important features. Although it introduces more computational complexity during compression and decompression.

3. PNG Compression:

- **Method:** Use the PNG format, which combines lossless compression algorithms
- **Impact:** Efficiently reduces storage size with minimal impact on runtime.

Storage and Runtime Costs:

- **Typical Storage Costs:**
 - **Uncompressed Image:** Each 100,000 x 100,000 pixel image requires approximately 10 GB of storage (assuming 1 byte per pixel).
 - **Compressed Image:** Compression techniques can reduce storage requirements by an order of magnitude. For example, block-based compression or PNG might reduce the size to 1-2 GB.
- **Runtime Costs:**
 - Compression and decompression add computational overhead, which can be minimized with efficient algorithms.
 - The actual runtime for processing a 100,000 x 100,000 pixel image using optimized techniques (like the NumPy-based approach) is manageable

6. Tools and Techniques Used to Solve the Challenge

Tools Used:

- **Python:** The primary programming language used for my implementation.
- **NumPy:** A library for numerical computations, used for efficient array operations and processing large images.

LLM Techniques:

- **Large Language Models (LLMs):** Used for initial drafting and explanation of concepts, ensuring clarity and correctness in the description of algorithms and processes.

Steps to Solve the Challenge:

- 1. Understanding Requirements:**
 - Analyze the requirements to determine the need for efficient storage and processing of high-resolution images.
- 2. Choosing Data Structures:**
 - Select appropriate data structures (e.g., 2D lists, NumPy arrays) for representing images to balance storage efficiency and processing speed.
- 3. Generating Simulated Images:**
 - Implement functions to generate realistic simulated images for testing purposes.
- 4. Cancer Detection Algorithm:**
 - Develop and optimize the cancer detection algorithm using NumPy for efficient processing.
- 5. Testing and Verification:**
 - Test the implementation with simulated images to ensure accuracy and performance.